

Martin Elsman, DIKU

Seminar at Chalmers University of Technology — December 18–19, 2018

- The region-based memory model
- A type-and-effect system for region-based memory management
- Region- and effect-polymorphism
- Region inference and arrow effects
- Combining Region-inference and garbage collection
- Exercises


Region-based Memory Management

QUESTION:

Can the Algol stack discipline be applied to languages with dynamic data structures and higher-order functions?

IDEAS:

- Organize the heap as a stack of regions.
- At runtime, allocate all values in regions.
- Perform *region inference*: Insert allocation and deallocation directives in the program code at compile time.



A Type-and-effect System for Region-annotated Expressions

REGION-ANNOTATED TERMS:

The grammars for *places* (p), *values* (v), and *expressions* (e) are as follows:

$$\begin{aligned} p &::= \rho \mid \bullet \\ v &::= d \textbf{in } p \mid \lambda x. e \textbf{in } p \\ e &::= v \mid x \mid d \textbf{at } p \mid e_1 e_2 \mid \lambda x. e \textbf{at } p \\ &\quad \mid \textbf{letregion } \rho \textbf{ in } e \end{aligned}$$

- A place of the form \bullet denotes a non-existing region. Thus, access to a value stored in a place \bullet signifies a reference to deallocated memory.
- Each value resides in a particular region, denoted by the “**in** p ” part of the value.
- Each value-creating expression, such as $d \textbf{at } \rho$, is annotated with the region in which the value goes at runtime.

A Small-step Contextual Dynamic Semantics

EVALUATION CONTEXTS:

$$E ::= [\cdot] \mid E e \mid v E \mid \mathbf{letregion} \ \rho \ \mathbf{in} \ E$$

NOTICE:

- Evaluation is allowed under **letregion**-constructs to model evaluation in the presence of non-empty region stacks.

- Reduction rules are of the form $e \rightsquigarrow e'$, which reads:

“The expression e reduces in one step to the expression e' .”

- For simplicity, we leave out a potential recursive construct:

$$e ::= \dots \mid \mathbf{fix} \ f(x) \ \mathbf{at} \ p = e$$

Reduction Rules

ALLOCATION

$$e \rightsquigarrow e'$$

$$d \text{ at } \rho \rightsquigarrow d \text{ in } \rho \qquad \lambda x.e \text{ at } \rho \rightsquigarrow \lambda x.e \text{ in } \rho$$

DEALLOCATION:

$$\text{letregion } \rho \text{ in } v \rightsquigarrow v[\bullet/\rho]$$

FUNCTION APPLICATION:

$$(\lambda x.e \text{ in } \rho) v \rightsquigarrow e[v/x]$$

CONTEXT:

$$\frac{e \rightsquigarrow e' \quad E \neq [\cdot]}{E[e] \rightsquigarrow E[e']} \quad (3)$$

Evaluation

We define *evaluation* as the least relation formed by the reflexive transitive closure of the reduction relation \rightsquigarrow :

EVALUATION

$$e \rightsquigarrow^* e'$$

$$\frac{e \rightsquigarrow e'}{e \rightsquigarrow^* e'} \quad (4)$$

$$\frac{}{e \rightsquigarrow^* e} \quad (5)$$

$$\frac{e_1 \rightsquigarrow^* e_2 \quad e_2 \rightsquigarrow^* e_3}{e_1 \rightsquigarrow^* e_3} \quad (6)$$

We further define $e \Uparrow$ to mean that there exists an infinite sequence,
 $e \rightsquigarrow e_1 \rightsquigarrow e_2 \rightsquigarrow \dots$.

A Region Type System

PURPOSE: provide a type system with the guarantee

“Well-typed programs do not get stuck.”

Types and type-and-places are defined by the grammars:

$$\mu ::= (\tau, \rho) \quad \text{(Type-and-places)}$$
$$\tau ::= \text{int} \mid \mu_1 \xrightarrow{\varphi} \mu_2 \quad \text{(Types)}$$

A *type environment* (Γ) maps program variables to type and places.

Type judgments $\Gamma \vdash e : \mu, \varphi$ are read:

“In the type environment Γ , the expression e has type and place μ and effect φ .”

Region Typing Rules

VALUES

$$\Gamma \vdash e : \mu, \varphi$$

$$\frac{}{\Gamma \vdash d \textbf{ in } \rho : (\text{int}, \rho), \emptyset} \quad (7) \qquad \frac{\{x : \mu_1\} \vdash e : \mu_2, \varphi}{\Gamma \vdash \lambda x. e \textbf{ in } \rho : (\mu_1 \xrightarrow{\varphi} \mu_2, \rho), \emptyset} \quad (8)$$

EXPRESSIONS:

$$\frac{}{\Gamma \vdash d \textbf{ at } \rho : (\text{int}, \rho), \{\rho\}} \quad (9) \qquad \frac{\Gamma + \{x : \mu_1\} \vdash e : \mu_2, \varphi}{\Gamma \vdash \lambda x. e \textbf{ at } \rho : (\mu_1 \xrightarrow{\varphi} \mu_2, \rho), \{\rho\}} \quad (10)$$

$$\frac{\Gamma(x) = \mu}{\Gamma \vdash x : \mu, \emptyset} \quad (11) \qquad \frac{\Gamma \vdash e_1 : (\mu' \xrightarrow{\varphi_0} \mu, \rho), \varphi_1 \quad \Gamma \vdash e_2 : \mu', \varphi_2}{\Gamma \vdash e_1 e_2 : \mu, \varphi_0 \cup \varphi_1 \cup \varphi_2 \cup \{\rho\}} \quad (12)$$

$$\frac{\Gamma \vdash e : \mu, \varphi \quad \varphi' \supseteq \varphi}{\Gamma \vdash e : \mu, \varphi'} \quad (13) \qquad \frac{\Gamma \vdash e : \mu, \varphi \quad \rho \notin \text{frv}(\Gamma, \mu)}{\Gamma \vdash \textbf{letregion } \rho \textbf{ in } e : \mu, \varphi \setminus \{\rho\}} \quad (14)$$

Properties of the Region Type System

LEMMA (TYPE PRESERVATION)

If $\vdash e : \mu, \varphi$ and $e \rightsquigarrow e'$ then $\vdash e' : \mu, \varphi$.

LEMMA (PROGRESS)

If $\vdash e : \mu, \varphi$ then e is a value or $e \rightsquigarrow e'$ for some e' .

The progress lemma implies that a well-typed program cannot apply a non-function to some argument or access values in regions that are deallocated.

THEOREM (TYPE SOUNDNESS)

If $\vdash e : \mu, \varphi$ then either $e \Uparrow$ or there exists some v such that $e \rightsquigarrow^ v$ and $\vdash v : \mu, \varphi$.*

Region and Effect Polymorphism

The type system extends naturally to support polymorphism in region variables and so-called *effect variables*, ranged over by ϵ .

- Effects (φ) are now sets of effect variables and region variables.
- Using effect variables, the type system can track higher-order programming with effects:
- The apply function $\lambda f. \lambda x. f \ x$ can be given the type

$$\forall \alpha \beta \rho \rho' \epsilon. ((\alpha, \rho) \xrightarrow{\{\epsilon\}} (\beta, \rho')) \xrightarrow{\emptyset} (\alpha, \rho) \xrightarrow{\{\epsilon\}} (\beta, \rho')$$

Region Inference

- Region Inference can be formulated both as a *constraint-based analysis* and as a *unification-based type inference algorithm*.
- Idea: identify arrow effects using effect variables.
- Arrow types take the form $\mu \xrightarrow{\epsilon.\varphi} \mu'$.

CONSTRAINT-BASES ANALYSIS (BIRKEDAL, TOFTE, JOURNAL OF THE.COMP.SC. '01):

- Whenever a type $\mu \xrightarrow{\epsilon.\varphi} \mu'$ is created, enforce the constraint $\epsilon \supseteq \varphi$.

UNIFICATION-BASED INFERENCE (TOFTE, BIRKEDAL, TOPLAS '98):

- When unifying two arrow effects $\epsilon.\varphi$ and $\epsilon'.\varphi'$, create a unifier:

$$S = \{\epsilon \mapsto \epsilon.\varphi'', \epsilon' \mapsto \epsilon.\varphi''\} \quad \text{where} \quad \varphi'' = S(\varphi) \cup S(\varphi')$$

Region Inference, Termination

Enforce consistency restrictions on introduced arrow effects.

A set of arrow effects Φ (called an *effect basis*) is said to be consistent if

- 1. It is functional:** For all $\epsilon.\varphi \in \Phi$ and $\epsilon'.\varphi' \in \Phi$, if $\epsilon = \epsilon'$ then $\varphi = \varphi'$.
- 2. It is closed:** For all $\epsilon.\varphi \in \Phi$ and $\epsilon' \in \varphi$, there exists φ' such that $\epsilon'.\varphi' \in \Phi$.
- 3. It is transitive:** For all $\epsilon.\varphi \in \Phi$ and $\epsilon'.\varphi' \in \Phi$, if $\epsilon' \in \varphi$ then $\varphi' \subseteq \varphi$.

Intuitively, a consistent basis form a set of DAGs.

A *contraction* (a substitution) is defined so that it is known to “shrink a basis”.

Unification is shown to generate contractions, which, in essence results in termination of region inference.

Garbage Collecting Regions

WHY COMBINE REGION INFERENCE AND GC?

- For non-regionized programs, adding GC reduces memory usage.
- From GC's point of view:
 - In general, region inference reduces the number of GC invocations.
 - Region-based memory management supports “almost tag-free” garbage collection.

A CHALLENGE:

- The Tofte-Talpin region typing rules permit dangling pointers!

Dangling Pointer Example

Consider the expression

$$\begin{aligned} e \quad &\equiv \quad \textbf{letregion } \rho \\ &\quad \textbf{in } (\lambda y. (\lambda x. (\lambda z. (1 \textbf{ at } \rho_1) \textbf{ at } \rho_1) y \textbf{ at } \rho_1) \textbf{ at } \rho_1) \\ &\quad \quad (3 \textbf{ at } \rho) \\ &\quad \textbf{end} \end{aligned}$$

From the typing rules, we have

$$\vdash e : ((\text{int}, \rho_1) \xrightarrow{\{\rho_1\}} (\text{int}, \rho_1), \rho_1), \{\rho_1\}$$

We also have (using **five** reduction steps)

$$e \rightsquigarrow^* \lambda x. (\lambda z. (1 \textbf{ at } \rho_1) \textbf{ at } \rho_1) (3 \textbf{ in } \rho) \textbf{ in } \rho_1$$

Problem: ρ , which is in the type of y , is not in the type of $\lambda x. (...)$

Disallowing Dangling Pointers

THE PROBLEM WITH THE TOFTE-TALPIN TYPING RULES:

- Values stored in function closures are not required to be contained in regions mentioned in the type of the function.

$$\frac{\Gamma + \{x : \mu_1\} \vdash e : \mu_2, \varphi}{\Gamma \vdash \lambda x.e \textbf{ at } \rho : (\mu_1 \xrightarrow{\varphi} \mu_2, \rho), \{\rho\}} \quad (15)$$

THE SOLUTION:

- Enforce region variables in the type of free variables of a function to appear in the type of the function.

$$\frac{\Gamma + \{x : \mu_1\} \vdash e : \mu_2, \varphi \quad \forall y \in \text{fv}(\lambda x.e). \text{frv}(\Gamma(y)) \subseteq \text{frv}(\mu_1 \xrightarrow{\varphi} \mu_2, \rho)}{\Gamma \vdash \lambda x.e \textbf{ at } \rho : (\mu_1 \xrightarrow{\varphi} \mu_2, \rho), \{\rho\}} \quad (16)$$

- The restriction has little impact on memory usage in practice (Elsman, TLDI'03).

Cheney's Algorithm for Regions

Extend Cheney's copying garbage collection algorithm to work with regions (Hallenberg, Elsman, Tofte, PLDI'02).

- Perform a Cheney copying collection on each region on the region stack.
- If a live value resides in a region r before a collection, the value must reside in the same region r after the collection.

HOW IT WORKS:

- All values reachable from the root set are evacuated into “to-space” region pages associated with the region.
- After a collection, all “from space” region pages are inserted into the free list of pages.

Example: Bootstrapping the MLKit

Compiling the MLKit with MLton produces an executable **Kit1** (takes 12min).

When running, **Kit1** uses the MLton runtime system.

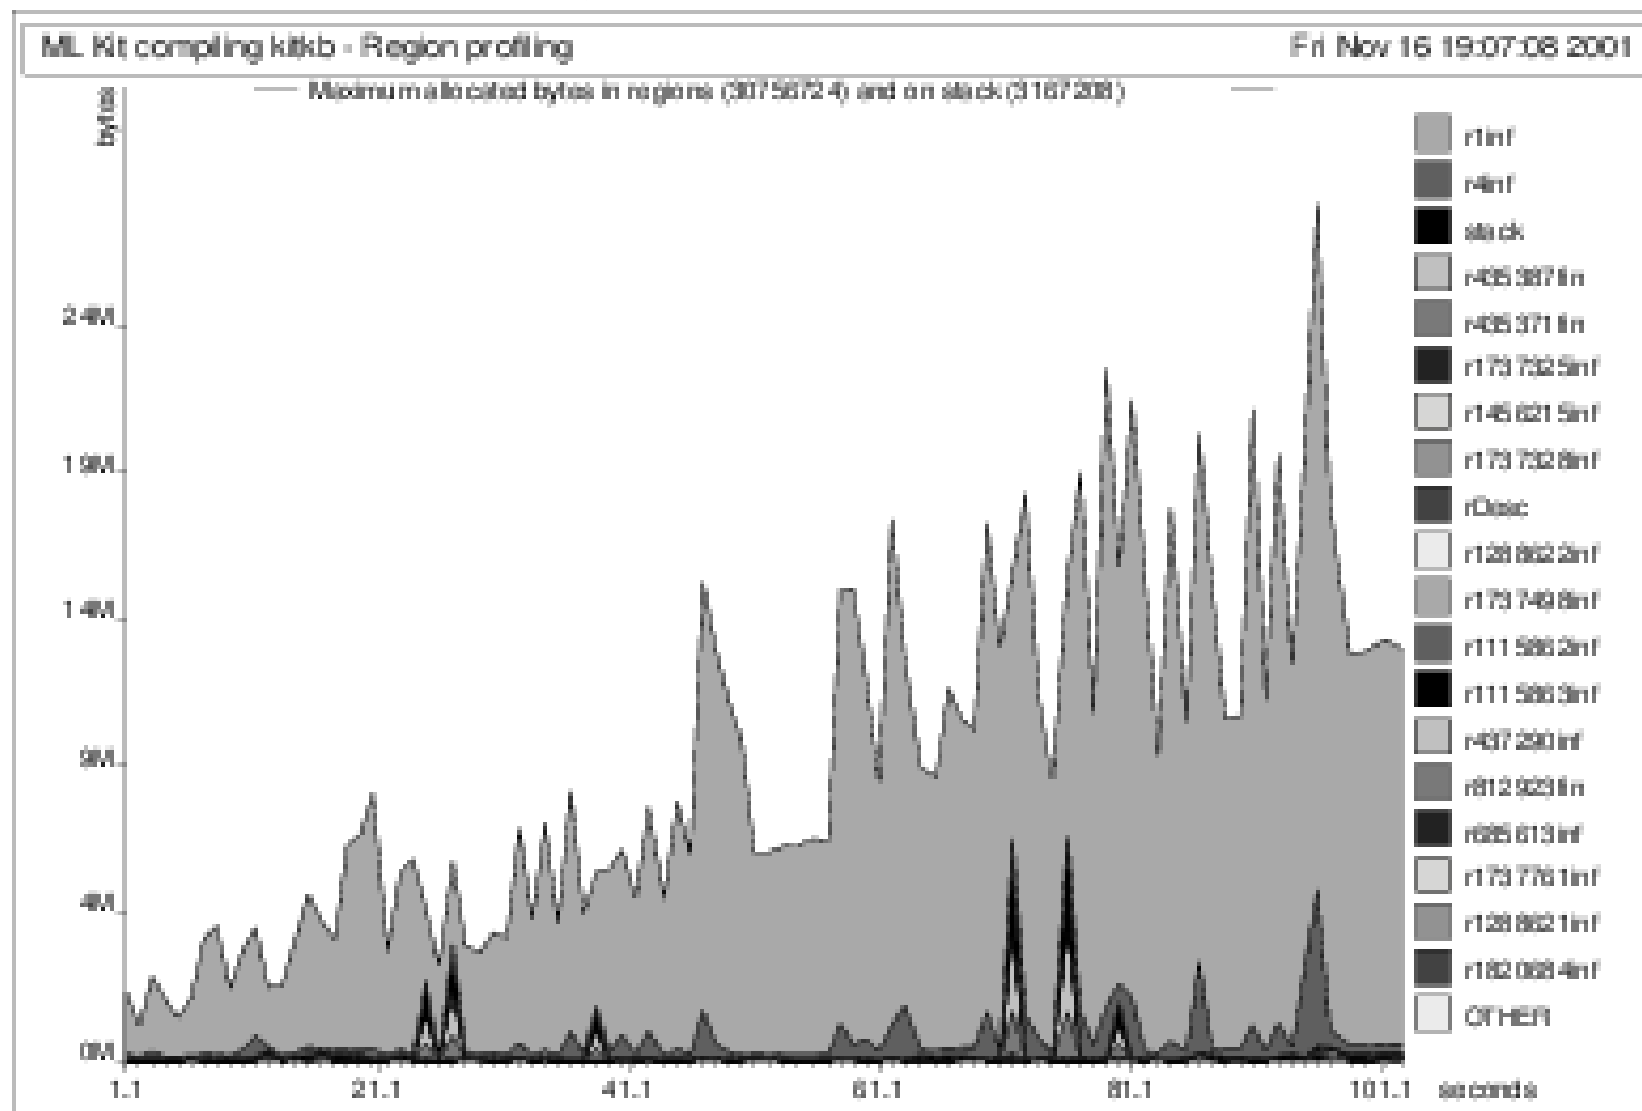
EXPERIMENTS:

- Using **Kit1** to compile the MLKit takes 3min and results in an executable **Kit2**, which uses a runtime system that combines regions and garbage collection.*
- Using **Kit2** to compile the MLKit takes 6min.
- But **Kit1** and **Kit2** are not whole-program compilers!

*The experiments were run on a MacBook Pro 2,7GHz Intel Core i7 laptop with 16Gb RAM.

Garbage Collecting Regions in Practice


EXAMPLE: A region profile of running the bootstrapped MLKit with the Knuth-Bendix test program as input:




Almost Tag-free Garbage Collection

Refine the region typing rules to disallow values of different types to reside in the same region.

- Runtime tags can then be moved from individual values to the region in which each value is stored.
- Dramatic savings in heap usage can be obtained, particularly for lists and tree-like data structures.



Untagged region
with tagged values



Tagged region
with untagged values

Possible Future Work


- Programming explicitly with regions
- Combining region inference and generational collection
- Thread support
- Region type systems for low-level languages (e.g., bytecode)

Exercises

① Clone the git-repository

<https://github.com/melsman/effects-seminar-public.git>

② Write a program that, in Dicken's Christmas Carol (available from the `texts/` folder), will find the number of sentences containing pairs of permuted words. Get the program to run in as little space as possible. Here is what you should aim for:



③ Prove the type soundness theorem on slide 3-9.